

THAITHE KALATHIL HARYSHWA NAIR

New Jersey, USA | +1 (551) 331-8898 | haryshwanair29@gmail.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

Security engineer focused on detection engineering, security automation, adversary emulation, and AI/LLM security. I build practical defensive systems, then break and test them to understand where they fail and how they can be improved.

EDUCATION

M.S. Cybersecurity, Stevens Institute of Technology — Hoboken, NJ Aug 2024 – May 2026
Minor: AI; Coursework: Advanced Cryptography, TCP/IP, Systems Programming, Privacy in a Networked World GPA: 3.43/ 4
B.Tech. Computer Science Engineering, Vellore Institute of Technology — Bhopal, India Aug 2020 – May 2024
Minor: Cybersecurity GPA: 7.92/ 10

TECHNICAL SKILLS & CERTIFICATIONS

Security: Detection Engineering, Threat Investigation, Incident Response, SIEM Monitoring, Log & PCAP Analysis, IAM, Adversary Emulation, AI/LLM Security, Threat Intelligence

Tools & Platforms: Splunk Enterprise Security, Wireshark, Docker, Linux, GCP, Git/GitHub, GitHub Actions, Ollama, SQLite, Nmap, Burp Suite, Metasploit, VirusTotal API

Development: Python, Bash, TypeScript, Security Automation

Certifications: Certified Ethical Hacker (CEH) | TryHackMe: Digital Forensics & Electronic Evidence, AI Security

TECHNICAL PROJECTS

Arbiter AI — Self-Hosted AI Security Alert Triage Engine — *Open-Source MVP, Active Development*

Python | Local LLMs (Ollama) | Docker | SQLite | GitHub Actions

- Open-sourced a triage engine for organizations with no dedicated SOC: a deterministic pre-filter scores severity, asset criticality, context and verdict history, escalating only ambiguous cases to a locally hosted LLM so no event data leaves the host. Built around asymmetric error costs — suppression requires strong benign evidence, while uncertain or model-failed cases escalate by policy.
- Built evaluation suites covering 296 labeled cases across 88 event types to benchmark local LLMs for security triage; discovered that small test sets produced unstable model rankings and used the larger evaluation to drive model-selection decisions.
- Found the model over-applying benign framing such as maintenance windows to disguised attacks, and moved attack-disguise checks into a non-suppressible guardrail layer the LLM cannot override. Unit tests and an end-to-end smoke run gate every commit in CI across Python 3.10–3.13.

TrustKit AI — Real-Time Voice & Vision Rental Fraud Copilot — *Collaborative AI Security Project*

Python | FastAPI | OpenCV | Vertex AI / Gemini | WebSockets | GCP

- Co-built a three-person fraud copilot that compares live virtual rental tours against listing data and external evidence, interrupting the viewer through a full-duplex voice agent the moment it detects a discrepancy.
- Contributed the backend frame extraction, metadata analysis and vision-processing modules feeding OpenCV forensic scoring and Vertex AI vision inference into the reasoning pipeline; continued development with the team past the hackathon.
- Adapted the verification pipeline when listing sites restricted automated access, combining available property data with submitted media and metadata to give the agent additional evidence for comparison.

SSH Honeypot — Adversary Emulation & Attack Telemetry — *Personal Security Lab*

Python | Docker | Kali Linux | Linux Security

- Built a medium-interaction SSH honeypot in isolated Docker containers capturing authentication attempts, commands and session activity, with no host filesystem exposure or outbound access.
- Attacked it from Kali Linux to generate controlled adversary behavior, then analyzed the captured command sequences from the defender's side.

WORK EXPERIENCE

Cybersecurity SOC Intern | Deloitte India Oct 2023 – Jan 2024

- Investigated 30–70 alerts per day inside a five-person remote SOC supporting banking and insurance clients on Splunk Enterprise Security environments processing 2,000+ security events daily.
- Correlated logs and validated alerts against client-specific playbooks, closing non-actionable noise and escalating confirmed or suspicious incidents to Tier 2 analysts.
- Documented findings to established SOC procedure, keeping event classification consistent across client workflows.

ACHIEVEMENTS & LEADERSHIP

- **Volunteer Lead, Google Developer Group Brooklyn** (Nov 2025 – Present) — One of four volunteer leads on a six-person organizing team across 8 events; coordinate 20+ volunteers, support speakers and sponsors, and run the chapter's LinkedIn communication.
- **Graduate Badminton Club Officer, Stevens Institute of Technology** (2024 – Present) — Weekly sessions, court logistics and equipment for a 250+ member club.
- **Black Belt & National Bronze Medalist, Shotokan Karate, India** (2008–2017) — Ten years of training and multiple state gold medals and a National bronze medal; later taught junior classes.